

# Adapters

Instead of the [Hermes primitives](#), existing applications use I/O libraries and middleware such as the C standard I/O library, MPI-IO, or HDF5. Since Hermes is intended to be a seamless I/O buffering solution, a set of *adapters* is provided in the form of LD\_PRELOAD-able shared libraries. These adapters perform the task of mapping between the application view of I/O primitives (e.g., files) and the Hermes primitives. An exemplary discussion of such mappings can be found in *Enosis: Bridging the Semantic Gap between File-based and Object-based Data Models* (<http://www.cs.iit.edu/~scs/assets/files/Enosis.pdf>) and *Syndesis: Mapping Objects to Files for a Unified Data Access System* (<http://www.cs.iit.edu/~scs/assets/files/Syndesis.pdf>).

Another important task of Hermes adapters is to collect I/O profiling that supports the detection of I/O patterns.

## Contents

### STDIO Adapter

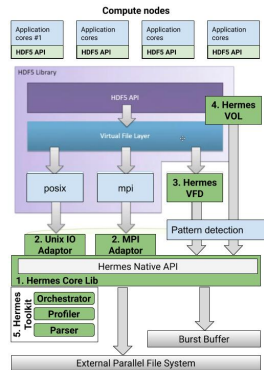
- Metadata structures
- Multi Thread support
- Questions/Considerations

### POSIX Adapter

### Unix Adapter

- Metadata structures
  - File table (FT)
  - File descriptor table (FDT)
  - Working set (WS)
  - Working set statistics (WSS)
- Supported APIs
  - open
  - write
  - read
  - lseek
  - posix\_fallocate
  - close
- System calls in H5FDsec2.c

### References



Hermes Adapters

# STDIO Adapter

The STDIO adapter assumes a paged view of a logical file, a byte sequence. The paged view uses a fixed page size (e.g., 1 MB) with which the logical file is partitioned within the adapter. The STDIO adapter will hint the Hermes library to use the page size as the buffer size configuration within the system. There are three types of APIs within STDIO.

- **Metadata APIs**
  - FILE access APIs
    - fclose
    - fopen
    - fopen64
    - fdopen
    - freopen
    - freopen64
  - FILE position APIs
    - fseek
    - fseeko
    - fseeko64
    - fsetpos
    - fsetpos64
    - rewind

- Bulk Data APIs
  - fread
  - fwrite
  - fflush
  - fgets
  - fputs
- Byte Data APIs
  - fgetc
  - putw
  - getc
  - getw
  - putc

The data structure that can maintain data and metadata info is local to the process. This is because FILE structure is in the scope of the process and needs locks using `flockfile`. Additionally, STDIO APIs do not support **shared file access** across multiple processes. Hence, we need to map user's FILE\* to our internal stat structure. The description of the structure is

```
struct FileID {
    dev_t dev_id_; /* ID of device containing file */
    ino_t inode_num_; /* inode number */
};
```


```
struct AdapterStat {
    BucketID    st_bkid; /* bucket associated with the file */
    i32         ref_count; /* # of time process opens a file */
    mode_t      st_mode; /* protection */
    uid_t       st_uid; /* user ID of owner */
    gid_t       st_gid; /* group ID of owner */
    off_t       st_size; /* total size, in bytes */
    off_t       st_ptr; /* Current ptr of FILE */
    blksize_t   st_blksize; /* blocksize for blob within bucket */
    time_t      st_atime; /* time of last access */
    time_t      st_mtime; /* time of last modification */
    time_t      st_ctime; /* time of last status change */
};
```

The proposed mapping of STDIO to Hermes constructs is as follows


```
FILE -> Bucket
FILE Pages -> Blobs
```

Where each file is a bucket (linked through bucket id). Each File contains pages where each page is referred through a key (corresponds to blob id). An illustration of the idea is presented below.


`fopen("/home/file1.dat", "w")`




`fwrite(write_buf, sizeof(char), 1024, fh)`




`fseek(fh, 384, SEEK_SET)`



`fwrite(write_buf, sizeof(char), 512, fh)`



`fwrite(write_buf, sizeof(char), 1024, fh)`





## Metadata structures

## Multi Thread support

In most cases, we assume we don't need multi-threading support as we will be employing process-level parallelism using MPI. However, if we are supporting Multi-Threading in the adapter, each critical path would be locked using `flockfile` on the in-memory file.

## Questions/Considerations

---

- STDIO does not support inter-process file sharing (only POSIX does).
- The Byte data APIs described above, perform I/O in byte granularity. In such a case the adapter performs blob updates in bytes.
- We need to either aggregate these requests in Adapter (?) or
- Optimize this special case within Hermes through hints.


## POSIX Adapter

## Unix Adapter


The Unix adapter assumes a paged view of a logical file, a byte vector.

What's the relationship between the page size and the pre-defined buffer sizes in buffer pools?


fopen("/home/file1.dat", "w")




fwrite(write\_buf, sizeof(char), 1024, fh)




fseek(fh, 384, SEEK\_SET)

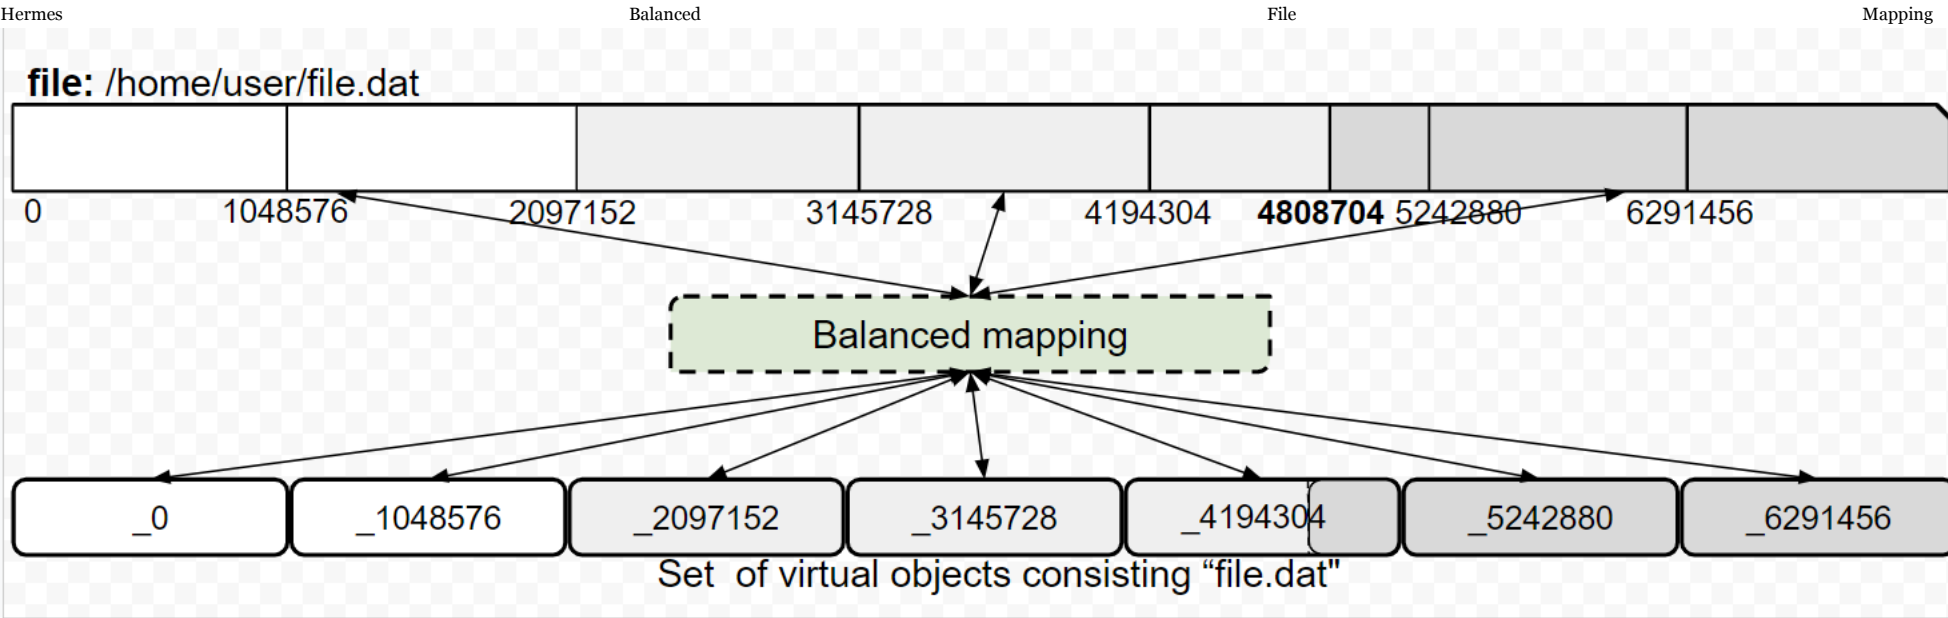


fwrite(write\_buf, sizeof(char), 512, fh)



fwrite(write\_buf, sizeof(char), 1024, fh)





This adapter supports a subset of `posix` calls (see below) related to file I/O. It also supports `poor-man's parallel I/O` to multiple independent files (MIF).

## Metadata structures

We call them tables, but they are maps

### File table (FT)

The **File ID** is obtained from `fstat`. The first two fields of the returned `struct stat` uniquely identify a file.

File ID := (st_dev, st_ino)
File ID -> (Bucket ID, Page Size, Reference Count)

Do we need to track EOA/EOF?

### File descriptor table (FDT)

File Desc. -> (File ID, Cursor, Append Mode)
--

### Working set (WS)

The set of page keys currently present in file buckets.

File Page Key := (File ID, Page Key)
--------------------------------------

### Working set statistics (WSS)

```
File Page Key -> (read count, write count)
```

**WARNING:** Generally, the key sets of WS and WSS will be different. The former is a subset of the latter.

## Supported APIs

### open

```
int open(const char *pathname, int flags);  
int open(const char *pathname, int flags, mode_t mode);
```

- Open/create the file, obtain new file descriptor
- Compute File ID
- IF File ID in FT:
  - Increment Reference Count
- ELSE:
  - Call flock and place an exclusive lock on the file
  - Acquire a Hermes bucket
  - Create a new entry in FT
  - Release the lock
- Add entry to FDT
- RETURN file descriptor

### write

```
ssize_t write(int fd, const void *buf, size_t count);
```

- Compute File Page Key Range
- IF File Desc. IN Append Mode:
  - Move Cursor to end of file
- FORALL File Page Key IN File Page Key Range:
  - IF writing full page:
    - Call hermes::put on file bucket w/ File Page Key and user buffer
  - ELSE:
    - IF File Page Key in WS:
      - Call hermes::get to retrieve buffer associated with File Page Key
    - ELSE:
      - Call client\_read to retrieve the data
    - Merge w/ user buffer
    - Call hermes::put
- UNLESS File Page Key in WSS:
  - Add (File Page Key, (0,0)) to WSS
- Increment write count for File Page Key
- Advance Cursor for this File Desc.
- RETURN the number of bytes written

### read

```
ssize_t read(int fd, void *buf, size_t count);
```

- Compute File Page Key Range
- FORALL File Page Key in Page Key Range:
  - IF File Page Key in WS:
    - hermes::get the data
  - ELSE:
    - Call client\_read to retrieve the data
    - hermes::put the data (optional?)
- IF reading full page:
  - Copy data to user buffer
- ELSE:
  - Merge data into user buffer



- UNLESS File Page Key in WSS:
  - Add (File Page Key, (0,0)) to WSS
  - Increment read count for File Page Key
- Advance Cursor for this File Desc.
- RETURN the number of bytes read

## lseek

```
off_t lseek(int fd, off_t offset, int whence);
```

- Call lseek
- Update the Cursor for File Desc. fd
- RETURN the new offset

## posix\_fallocate

```
int posix_fallocate(int fd, off_t offset, off_t len);
```

## close

```
int close(int fd);
```

- Call client\_close
- Remove the corresponding entry from FDT
- Decrement the Reference Count
- IF Reference Count == 0:
  - Call client\_flock and place an exclusive lock on the file
  - Release the Hermes bucket
  - Remove the entry from FT
  - Release the lock

## System calls in H5FDsec2.c

- open
- close
- read, pread
- write, pwrite
- lseek
- stat
- ftruncate
- flock
- fallocate (?)

## References

- [GOTCHA is a library for wrapping function calls in shared libraries \(https://github.com/LLNL/GOTCHA\)](https://github.com/LLNL/GOTCHA)

Retrieved from "<https://hermes.page/index.php?title=Adapters&oldid=834>"